

```
In [1]: 1 # Importing necessary libraries
2
3 import pandas as pd
4 import numpy as np
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7 from sklearn.model_selection import train_test_split
8 from sklearn.linear_model import LinearRegression
9 from sklearn.metrics import mean_squared_error, r2_score
10
11 # Load the dataset
12 med_ins = pd.read_csv(r'C:\Users\Rohit kumar\Downloads\pythond\medical_insurance.csv')
13
14 # Display the first few rows of the dataset
15 print("First few rows of the dataset:")
16 print(med_ins.head())
```

First few rows of the dataset:

	age	sex	bmi	children	smoker	region	charges
0	19	0	27.900	0	0	1	16884.92400
1	18	1	33.770	1	1	2	1725.55230
2	28	1	33.000	3	1	2	4449.46200
3	33	1	22.705	0	1	3	21984.47061
4	32	1	28.880	0	1	3	3866.85520

```
In [2]: 1 # Statistical summary of the dataset
2 print("Statistical summary of the dataset:")
3 print(med_ins.describe())
4
5 # Check for any missing values
6 print("\nNumber of missing values in each column:")
7 print(med_ins.isnull().sum())
8
```

Statistical summary of the dataset:

	age	sex	bmi	children	smoker
count	2772.000000	2772.000000	2772.000000	2772.000000	2772.000000
mean	39.109668	0.507215	30.701349	1.101732	0.796537
std	14.081459	0.500038	6.129449	1.214806	0.402647
min	18.000000	0.000000	15.960000	0.000000	0.000000
25%	26.000000	0.000000	26.220000	0.000000	1.000000
50%	39.000000	1.000000	30.447500	1.000000	1.000000
75%	51.000000	1.000000	34.770000	2.000000	1.000000
max	64.000000	1.000000	53.130000	5.000000	1.000000

	region	charges
count	2772.000000	2772.000000
mean	2.467532	13261.369959
std	1.103467	12151.768945
min	1.000000	1121.873900
25%	2.000000	4687.797000
50%	2.000000	9333.014350
75%	3.000000	16577.779500
max	4.000000	63770.428010

Number of missing values in each column:

```
age      0
sex      0
bmi      0
children 0
smoker   0
region   0
charges  0
dtype: int64
```

```
In [3]: 1 #check the data type
2 print("Data type")
3 print med_ins.dtypes
4
```


```
Data type
age          int64
sex          int64
bmi          float64
children     int64
smoker       int64
region       int64
charges      float64
dtype: object
```

```
In [4]: 1 X = med_ins.iloc[:, [0,1,2,3,4,5]].values
2 Y = med_ins.iloc[:, -1].values
3
4 print(X)
```

```
[[19.    0.    27.9    0.    0.    1. ]
 [18.    1.    33.77   1.    1.    2. ]
 [28.    1.    33.     3.    1.    2. ]
 ...
 [19.    1.    26.03   1.    0.    3. ]
 [23.    1.    18.715  0.    1.    3. ]
 [54.    1.    31.6    0.    1.    1. ]]
```

```
In [31]: 1 # Visualize the relationships between the independent variables and the target variable
2 sns.pairplot(med_ins, x_vars=['age', 'sex', 'bmi', 'children', 'smoker', 'region'], y_var='charges',
3             plt.show())
4
```

C:\Users\Rohit kumar\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning:  
The figure layout has changed to tight  
self.\_figure.tight\_layout(\*args, \*\*kwargs)



```
In [7]: 1 X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)
```

```
In [8]: 1 regressor = LinearRegression()
2 regressor.fit(X_train, Y_train)
```

```
Out[8]: LinearRegression
LinearRegression()
```

```
In [9]: 1 Y_pred = regressor.predict(X_test)
2 np.set_printoptions(precision=2)
3 print(np.concatenate((Y_pred.reshape(len(Y_pred),1), Y_test.reshape(len(Y_test),1)),1))

[[ 1731.92  2221.56]
 [   27.54 21595.38]
 [ 7754.66  5327.4 ]
 ...
 [ 9721.78  9301.89]
 [13086.02  7650.77]
 [ 4689.88  6640.54]]
```

```
In [29]: 1 # Define the input features for prediction
2 input_data = [[25, 1, 30, 3, 0, 1]] # age=25, sex=male, bmi=30, children=3, smoker=1
3
4 # Make prediction
5 predicted_charge = regressor.predict(input_data)[0]
6 predicted_charge_round= round(predicted_charge,2)
7 print("Predicted medical charge for the given scenario:", "Rs",predicted_charge_round)

Predicted medical charge for the given scenario: Rs 29054.92
```


```
In [19]: 1 print(regressor.intercept_, regressor.coef_, regressor.score(X, Y))

11329.413044173816 [ 253.28 -145.14  317.56  545.44 -24052.04  375.58] 0.7507
060330466141
```

```
In [25]: 1 result = r2_score(Y_test, Y_pred)
2 R2 = round(result, 2)
3 R2_percent = R2 * 100
4 print("The R^2 is:", R2_percent, "percent")

The R^2 is: 75.0 percent
```

```
In [32]: 1 # Make predictions on the test set
2 Y_pred = regressor.predict(X_test)
3
4 # Scatter plot of actual charges vs. predicted charges
5 plt.figure(figsize=(8, 6))
6 plt.scatter(Y_test, Y_pred, color='blue', alpha=0.5)
7 plt.title('Actual Charges vs. Predicted Charges')
8 plt.xlabel('Actual Charges')
9 plt.ylabel('Predicted Charges')
10 plt.show()
11
```



```
In [ ]: 1
```